

We introduce a new language representation model called BERT, which stands for Bidirectional Encoder

The Masked Language Model (MLM): In order to train a deep bidirectional representation, we randomly

Next Sentence Prediction (NSP): Many important downstream tasks such as Question Answering and

BERT architecture is a multi-layer bidirectional Transformer encoder based on the original Transformer

On the GLUE benchmark, BERT achieves an overall score of 80.5%, obtaining substantial improvement

Fine-tuning BERT on SQuAD v1.1 question answering achieves 93.2 F1 score, outperforming human