

The dominant sequence transduction models are based on complex recurrent or convolutional neural

An attention function can be described as mapping a query and a set of key-value pairs to an output. T

Scaled Dot-Product Attention: We compute the attention function on a set of queries simultaneously, p

Multi-Head Attention: Instead of performing a single attention function with d_{model} -dimensional keys,

Positional Encoding: Since our model contains no recurrence and no convolution, in order for the model

On the WMT 2014 English-to-German translation task, the big transformer model achieves a state-of-t